

# Question 1

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: col_names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'Class']
df = pd.read_csv('iris.csv', names = col_names)
df
```

Out[2]:

	sepal-length	sepal-width	petal-length	petal-width	Class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...	...	...	...	...	...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 5 columns

```
In [3]: from math import sqrt
def euclidean_distance(row1, row2):
    distance = 0.0
    for i in range(len(row1)-1):
        distance += (row1[i] - row2[i])**2
    return sqrt(distance)

dataset = [[5.1,3.5,1.4,0.2],
[4.9,3.0,1.4,0.2],
[4.7,3.2,1.3,0.2],
[4.6,3.1,1.5,0.2],
[5.0,3.6,1.4,0.2],
[5.4,3.9,1.7,0.4],
[4.6,3.4,1.4,0.3],
[5.0,3.4,1.5,0.2],
[4.4,2.9,1.4,0.2],
[4.9,3.1,1.5,0.1]]
row0 = dataset[0]
for row in dataset:
    distance = euclidean_distance(row0, row)
    print(distance)
```

```
0.0
0.5385164807134502
0.509901951359278
0.648074069840786
0.1414213562373093
0.5830951894845303
0.5099019513592785
0.17320508075688762
0.9219544457292882
0.4582575694955836
```

```
In [4]: def get_neighbors(train, test_row, num_neighbors):
    distances = list()
    for train_row in train:
        dist = euclidean_distance(test_row, train_row)
        distances.append((train_row, dist))
    distances.sort(key=lambda tup: tup[1])
    neighbors = list()
    for i in range(num_neighbors):
        neighbors.append(distances[i][0])
    return neighbors

neighbors = get_neighbors(dataset, dataset[0], 3)
for neighbor in neighbors:
    print(neighbor)
```

```
[5.1, 3.5, 1.4, 0.2]
[5.0, 3.6, 1.4, 0.2]
[5.0, 3.4, 1.5, 0.2]
```

```
In [5]: def predict_classification(train, test_row, num_neighbors):
    neighbors = get_neighbors(train, test_row, num_neighbors)
    output_values = [row[-1] for row in neighbors]
    prediction = max(set(output_values), key=output_values.count)
    return prediction

prediction = predict_classification(dataset, dataset[0], 3)
print('Expected %d, Got %d.' % (dataset[0][-1], prediction))
```

```
Expected 0, Got 0.
```

## Question 2

```
In [6]: from sklearn.model_selection import train_test_split
X = df[['sepal-length', 'sepal-width', 'petal-length', 'petal-width']]
y = df[['Class']]
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.4)
```


```
In [111]: from sklearn import metrics
from sklearn.neighbors import KNeighborsClassifier
k_range = range(1,10,2)
scores=[]
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    scores.append(metrics.accuracy_score(y_test, y_pred))

print(scores)
```

```
[0.9666666666666667, 0.9833333333333333, 0.9833333333333333, 1.0, 1.0]
```

```
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:1
79: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:1
79: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:1
79: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:1
79: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:1
79: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
```

```
In [112]: plt.plot(k_range, scores)
plt.xlabel('K values')
plt.ylabel('Accuracy')
plt.show()
```



This graph shows that we are getting accuracy of 0.983 when k is between 3 to 5 but when we move ahead we found that k=7 gives the best accuracy so knee point must be floor then k must be 7

## Question 3

```
In [7]: from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import average_precision_score
from sklearn.naive_bayes import GaussianNB
df = pd.read_csv('wine.csv')
X = df.drop('Proline', axis=1)
y = df['Proline']

Accuracy = []
prec_mat = []
split_size = [0.5, 0.4, 0.3, 0.2, 0.1]
for i in range(0,5):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = split_size[i])
    gnb = GaussianNB()
    gnb.fit(X_train, y_train)

    y_pred = gnb.predict(X_test)

    Accuracy.append(metrics.accuracy_score(y_test, y_pred)*100)
    y_true = y_test.values
    y_true = y_true.reshape(-1,1)
    prec_mat.insert(i,metrics.classification_report(y_true, y_pred, digits = 3))

_warn_prf(average, modifier, msg_start, len(result))
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to
0.0 in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set t
o 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to
0.0 in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set t
o 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
... ..
```

```
In [8]: for i in range(0,5):
        print(prec_mat[i])
```

	precision	recall	f1-score	support
278	0.000	0.000	0.000	1
290	0.000	0.000	0.000	1
312	0.000	0.000	0.000	1
325	0.000	0.000	0.000	1
345	0.000	0.000	0.000	1
352	0.000	0.000	0.000	1
355	0.000	0.000	0.000	1
365	0.000	0.000	0.000	1
372	0.000	0.000	0.000	1
378	0.000	0.000	0.000	1
380	0.000	0.000	0.000	2
410	0.000	0.000	0.000	1
415	0.000	0.000	0.000	1
420	0.000	0.000	0.000	1
428	0.000	0.000	0.000	1
434	0.000	0.000	0.000	1
450	0.000	0.000	0.000	1
...	...	...	...	...

```
In [9]: plt.show(prec_mat[i])
```

# Question 4

```
In [10]: import pandas as pd
df = pd.read_csv('train.csv', delimiter=',')
df
```

Out[10]:

	Phraseld	Sentenceld	Phrase	Sentiment
0	1	1	A series of escapades demonstrating the adage ...	1
1	2	1	A series of escapades demonstrating the adage ...	2
2	3	1	A series	2
3	4	1	A	2
4	5	1	series	2
...	...	...	...	...
156055	156056	8544	Hearst 's	2
156056	156057	8544	forced avuncular chortles	1
156057	156058	8544	avuncular chortles	3
156058	156059	8544	avuncular	2
156059	156060	8544	chortles	2

156060 rows × 4 columns

```
In [11]: df.head()
```

```
Out[11]:
```

	Phraselid	Sentencelid	Phrase	Sentiment
0	1	1	A series of escapades demonstrating the adage ...	1
1	2	1	A series of escapades demonstrating the adage ...	2
2	3	1	A series	2
3	4	1	A	2
4	5	1	series	2

```
In [12]: from sklearn.feature_extraction.text import CountVectorizer
from nltk.tokenize import RegexpTokenizer
token = RegexpTokenizer(r'[a-zA-Z0-9]+')
cv = CountVectorizer(lowercase=True, stop_words='english', ngram_range = (1,1), tokenizer = token)
text_counts= cv.fit_transform(df['Phrase'])
```

```
In [13]: Accuracy = []
prec_mat = []
split_size = [0.5, 0.4, 0.3, 0.2, 0.1]
for i in range(0,5):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = split_size[i])
    gnb = GaussianNB()
    gnb.fit(X_train, y_train)

    y_pred = gnb.predict(X_test)

    Accuracy.append(metrics.accuracy_score(y_test, y_pred)*100)
    y_true = y_test.values
    y_true = y_true.reshape(-1,1)
    prec_mat.insert(i,metrics.classification_report(y_true, y_pred, digits = 3))
```

```
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to
0.0 in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Precision and F-score are ill-defined and being set t
o 0.0 in labels with no predicted samples. Use `zero_division` parameter to control
this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Nupur goel\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1245: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to
0.0 in labels with no true samples. Use `zero_division` parameter to control this b
ehavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [14]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn import metrics
split_size = [0.5, 0.4, 0.3, 0.2, 0.1]
for i in range(0,5):
    X_train, X_test, y_train, y_test = train_test_split(
        text_counts, df['Sentiment'], test_size = split_size[i], random_state=0)
    clf = MultinomialNB().fit(X_train, y_train)
    predicted= clf.predict(X_test)
    print("MultinomialNB Accuracy:",metrics.accuracy_score(y_test, predicted))
```

MultinomialNB Accuracy: 0.5956427015250545

MultinomialNB Accuracy: 0.6035979751377675

MultinomialNB Accuracy: 0.6095305224486308

MultinomialNB Accuracy: 0.6106946046392413

MultinomialNB Accuracy: 0.6144431628860695

In [ ]: